

CSE 617: Digital Image Processing

Local Feature Extraction

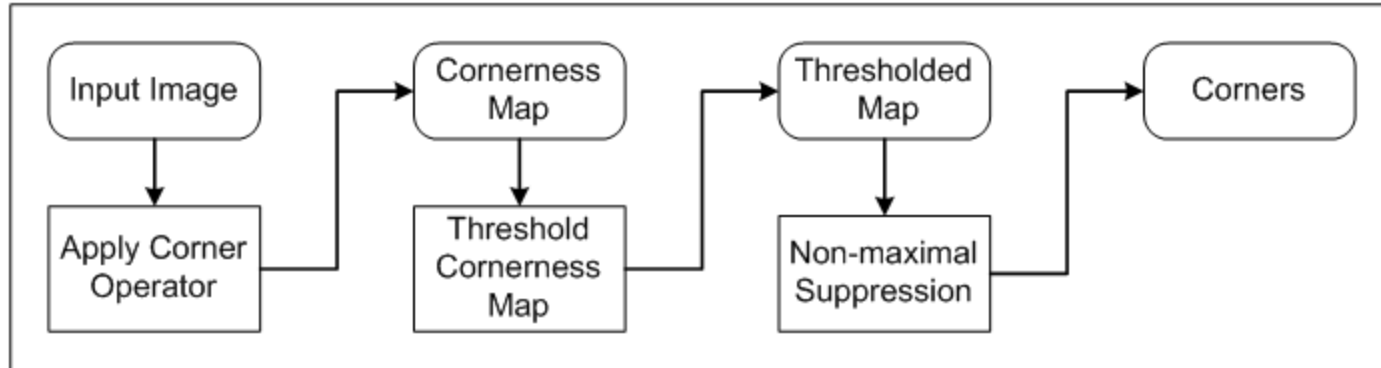
Seif Eldawlatly

Corner Detectors

- It is always useful to find pairs of corresponding points in two similar images
- This could be used in the analysis of moving images
- This could be done by comparing all possible pairs of pixels in the two images. However, it is computationally intensive
- This process might be simplified by comparing interest points only such as corners
- A corner can be defined as a pixel in its small neighborhood where two dominant and different edges meet

Moravec Operator

- General structure of corner detectors




<http://kiwi.cs.dal.ca/~dparks/CornerDetection/algorithms.htm>

- Moravec operator estimates the *cornerness* of a point by computing a measure of intensity variation in a given neighborhood w with a shift of u and v

$$V(x, y)_{w(u, v)} = \sum_{i=-1}^1 \sum_{j=-1}^1 (f(x+i, y+j) - f(x+u+i, y+v+j))^2$$

Moravec Operator

- It measures the cornerness by shifting a window around the considered pixel by 1 pixel in each of the 8 principal directions and calculating the corresponding V



<http://kiwi.cs.dal.ca/~dparks/CornerDetection/algorithms.htm>

Window B Centered at

$$V = \begin{bmatrix} 3 * 255^2 \\ 2 * 255^2 \\ 3 * 255^2 \\ 2 * 255^2 \\ 5 * 255^2 \\ 2 * 255^2 \\ 3 * 255^2 \\ 2 * 255^2 \end{bmatrix} \begin{matrix} A3 \\ A2 \\ A1 \\ A4 \\ A7 \\ A8 \\ A9 \\ A6 \end{matrix}$$

- The cornerness at a pixel (x, y)

$$C(x, y) = \min_{u, v} V(x, y)_{w(u, v)}$$


Moravec Operator

- Example (using a 3 x 3 window)

X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
X	X	0	0	0	0	0	0	0	0	0	1	1	1	X	X
X	X	0	0	0	0	0	1	1	0	0	1	2	1	X	X
X	X	0	0	0	0	0	2	1	0	0	1	1	1	X	X
X	X	0	0	0	0	0	0	0	0	0	0	0	0	X	X
X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X


Moravec Operator

- Why would Moravec operator work?



<http://kiwi.cs.dal.ca/~dparks/CornerDetection/algorithms.htm>

- By setting all points with corneriness below a threshold T to 0, corner points can be detected




Moravec Operator

- Finally, non-maximal suppression can be applied


Original Image




After Thresholding



**After Non-maximal
Suppression**



Final Output



Cornerness of any pixel is set to 0 if it is not larger than the cornerness of all 8-neighbors

Harris Corner Detector

- Improved upon Moravec operator

$$S_W(\Delta x, \Delta y) = \sum_{x_i \in W} \sum_{y_i \in W} (f(x_i, y_i) - f(x_i - \Delta x, y_i - \Delta y))^2$$

- Similar to the goal of Moravec operator, we try to find the minimum value of S_W
- This could be found analytically if the shifted image patch is approximated by the first-order Taylor expansion

$$f(x_i - \Delta x, y_i - \Delta y) \approx f(x_i, y_i) + \left[\frac{\partial f(x_i, y_i)}{\partial x}, \frac{\partial f(x_i, y_i)}{\partial y} \right] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

- Substituting in the expression for S_W

$$S_W(\Delta x, \Delta y) = \sum_{x_i \in W} \sum_{y_i \in W} \left(f(x_i, y_i) - f(x_i, y_i) - \left[\frac{\partial f(x_i, y_i)}{\partial x}, \frac{\partial f(x_i, y_i)}{\partial y} \right] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2$$

Harris Corner Detector

$$\begin{aligned} S_W(\Delta x, \Delta y) &= \sum_{x_i \in W} \sum_{y_i \in W} \left(f(x_i, y_i) - f(x_i, y_i) - \left[\frac{\partial f(x_i, y_i)}{\partial x}, \frac{\partial f(x_i, y_i)}{\partial y} \right] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2 \\ &= \sum_{x_i \in W} \sum_{y_i \in W} \left(- \left[\frac{\partial f(x_i, y_i)}{\partial x}, \frac{\partial f(x_i, y_i)}{\partial y} \right] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2 \\ &= \sum_{x_i \in W} \sum_{y_i \in W} \left(\left[\frac{\partial f(x_i, y_i)}{\partial x}, \frac{\partial f(x_i, y_i)}{\partial y} \right] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \right)^2 \\ &= \sum_{x_i \in W} \sum_{y_i \in W} [\Delta x, \Delta y] \left(\begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} \begin{bmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} \end{bmatrix} \right) \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \\ &= [\Delta x, \Delta y] \left(\sum_{x_i \in W} \sum_{y_i \in W} \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} \begin{bmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} \end{bmatrix} \right) \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \\ &= [\Delta x, \Delta y] A_W(x, y) \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}, \end{aligned}$$

Harris Corner Detector

- The goal now is to minimize

$$S_W(\Delta x, \Delta y) = [\Delta x, \Delta y] A_W(x, y) \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

- This is equivalent to finding the eigenvector of A_w corresponding to the minimum eigenvalue
- What are the eigenvectors and eigenvalues?

For any matrix L , the eigenvectors and eigenvalues are defined as

$$Lf = \lambda f \rightarrow f^T Lf = \lambda$$

f is an eigenvector of L

λ is the eigenvalue corresponding to f

- For A_w , the eigenvector is $\begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$ and the eigenvalue is the corresponding S_w

Harris Corner Detector

- Therefore, finding the eigenvalues of A_w would be sufficient to solve the minimization problem where

$$A(x, y) = \begin{bmatrix} \sum_{x_i \in W} \sum_{y_i \in W} \frac{\partial^2 f(x_i, y_i)}{\partial x^2} & \sum_{x_i \in W} \sum_{y_i \in W} \frac{\partial f(x_i, y_i)}{\partial x} \frac{\partial f(x_i, y_i)}{\partial y} \\ \sum_{x_i \in W} \sum_{y_i \in W} \frac{\partial f(x_i, y_i)}{\partial x} \frac{\partial f(x_i, y_i)}{\partial y} & \sum_{x_i \in W} \sum_{y_i \in W} \frac{\partial^2 f(x_i, y_i)}{\partial y^2} \end{bmatrix}$$

- Instead of computing the eigenvalues, Harris suggested using the following approximation

$$R(A) = \det(A) - \kappa \text{trace}^2(A)$$

where $\det(A)$ is the determinant of the local structure matrix A


$\text{trace}(A)$ is the trace of matrix A (sum of elements on the diagonal)

κ is a tunable parameter


Harris Corner Detector

- Algorithm
 1. Filter the image with a Gaussian filter
 2. Estimate intensity gradient in 2 perpendicular directions for each pixel
$$\frac{\partial f(x, y)}{\partial x} \text{ and } \frac{\partial f(x, y)}{\partial y}$$
 3. For each pixel and a given neighborhood window
 - Calculate the local structure matrix A
 - Evaluate the response function $R(A)$
 4. Set all pixels with response less than a threshold T to 0 and perform non-maximal suppression

Moravec Operator



Harris Detector



Scale Invariant Feature Transform (SIFT)

- References:

D.G. Lowe, “Distinctive Image features from scale-invariant keypoints,” International Journal on Computer Vision 60 (2), 91–110, 2004


Local Feature Extraction

- Image matching is a fundamental aspect of many problems in computer vision and image processing
- Applications include object or scene recognition, solving for 3D structure from multiple images, stereo correspondence, and motion tracking
- Example: Consider the problem of stitching multiple images to create a panoramic view



Local Feature Extraction

- Stitching two images can be easily achieved once pairs of matched points across the two images are found



- By mapping the correspondence, a panorama can be easily created



Local Feature Extraction

- This process can be achieved through two steps:

- Step 1: Identify keypoints in both images



- Step 2: Match keypoints to identify their correspondence



- Step 1 is needed to reduce the search space in Step 2

Local Feature Extraction


- Another application that can make use of keypoint detection is recognizing similar objects in different images



- In this application, objects (keypoints) should be identified independent of their scales
- Moravec operator or Harris corner detectors can be used for this application


Scale Invariant Detection

- An important approach that could be used for both images stitching or different scales recognition is scale invariant detection
- The goal of such method is to identify a keypoint region as the region whose brightness is different from the surrounding
- A very common method is to identify **Blobs** in which the circle of smallest radius that encloses a keypoint region is identified



Scale Invariant Detection

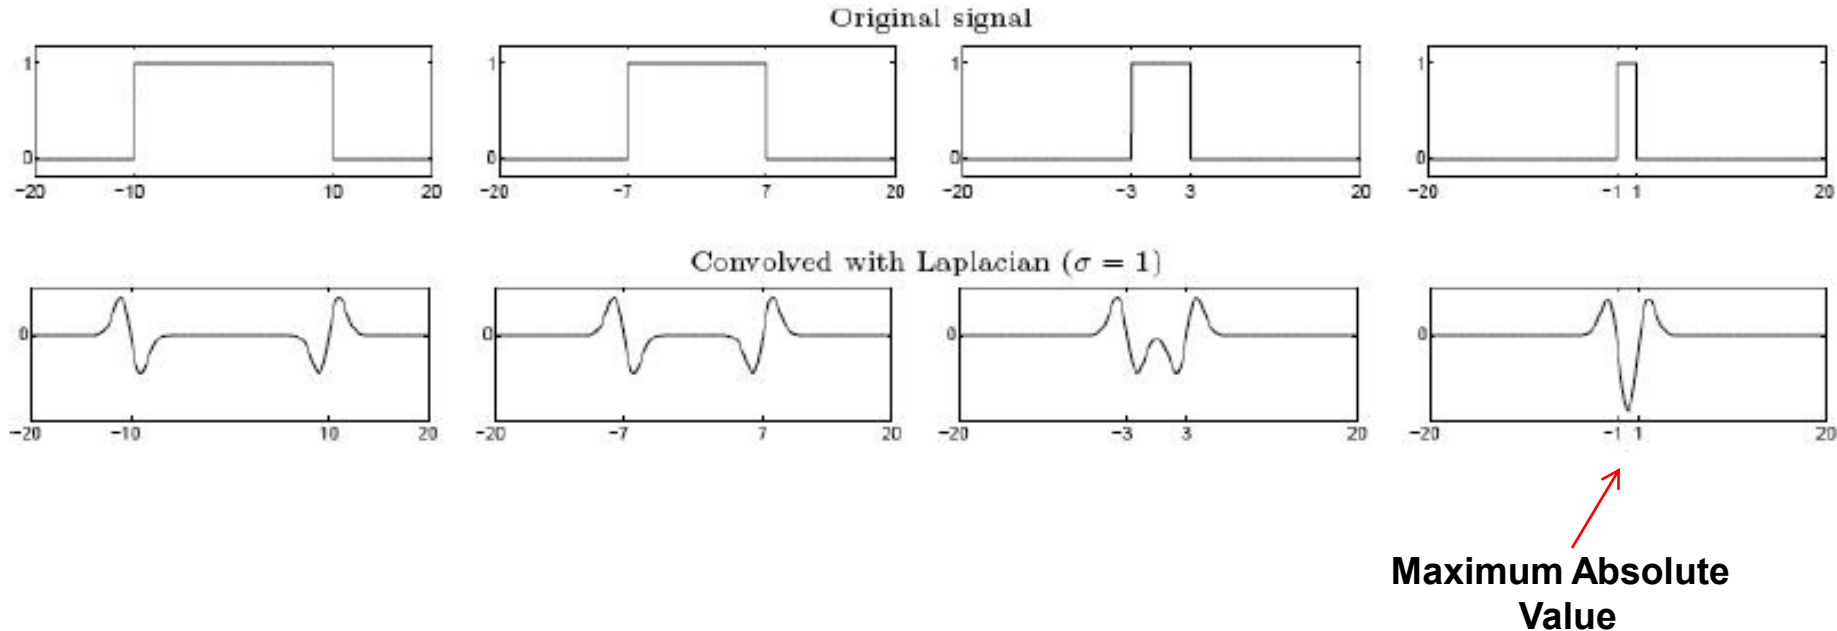
- Identifying blobs is based on using the Laplacian of Gaussian (LoG) method we previously encountered in edge detection



Source: S. Seitz

Scale Invariant Detection

- Applying the LoG mask to a blob gives the following response as the size of the blob increases



- If the scale of the LoG is matched to the blob size, the response of the LoG will be maximum at the center of the blob

Scale Invariant Detection

- Based on the previous observation, identifying a blob can be done by searching for the maximum absolute LoG response in:

- **Space**: Searching in local neighborhood
- **Scale**: Searching across different scales



Scale Invariant Feature Transform (SIFT)

- The SIFT algorithm can identify blobs and match them across different images
- SIFT Algorithm Steps:
 - Scale-space extrema detection
 - Orientation assignment
 - Generation of keypoint descriptors

SIFT: Scale-space Extrema Detection

- Keypoints in SIFT correspond to local extrema after applying the LoG mask across different scales
- The LoG mask represents a circularly symmetric operator for blob detection

$$LoG(x, y) = \frac{-1}{\pi\sigma^4} \left(1 - \frac{x^2 + y^2}{2\sigma^2} \right) e^{-\frac{x^2 + y^2}{2\sigma^2}}$$





SIFT: Scale-space Extrema Detection

- To normalize the values, the LoG is multiplied by σ^2

$$\sigma^2 LoG(x, y) = \frac{-1}{\pi\sigma^2} \left(1 - \frac{x^2 + y^2}{2\sigma^2} \right) e^{-\frac{x^2 + y^2}{2\sigma^2}}$$


- For a given keypoint, the corresponding best scale σ_{best} is determined as the one that maximizes the squared LoG response

- The corresponding radius is approximately $\sqrt{2}\sigma_{best}$




Identified Blob

Scale of Maximum
Squared LoG Response



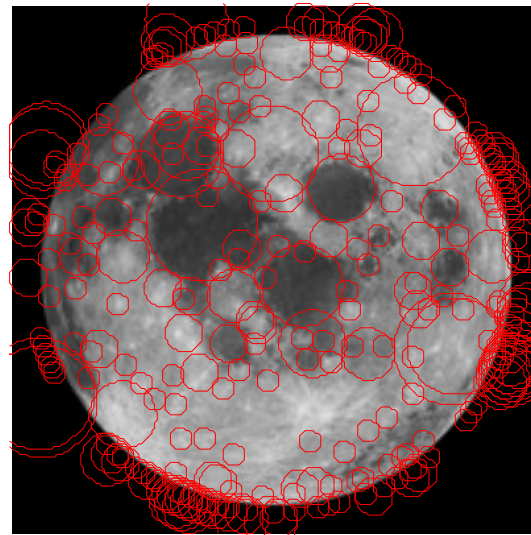

SIFT: Scale-space Extrema Detection

- Steps of the Scale-space Extrema Detection:
 - 1- Convolve the image with the normalized LoG for different values of σ
 - 2- Find the maxima of the squared LoG response in scale-space
- Maxima of the squared LoG images are detected by comparing a pixel (marked with X) to its 26 neighbors in 3x3 regions at the current and adjacent scales (marked with circles)




SIFT: Scale-space Extrema Detection

- Examples:



SIFT: Scale-space Extrema Detection

- Obtained extrema less than a predefined threshold t are usually ignored



SIFT: Orientation Assignment

- In order to stitch two images, blobs in both images are first identified and then compared

Image 1



Image 2



Blobs of Image 1



Blobs of Image 2




SIFT: Orientation Assignment

- Blobs corresponding to each other across the two images are then identified
- This is done in SIFT by first assigning a consistent orientation to each keypoint based on local image properties
- This is done as follows:
 - For each keypoint, convolve the image with a Gaussian filter with standard deviation of σ_{best}
 - Apply a first-derivative operator (for example Prewitt operator) to find the corresponding gradient image
 - Find the gradient orientation at each of the pixels in the 16 x 16 neighborhood around the keypoint

SIFT: Descriptors Generation

- For each blob, a descriptor is then constructed
- This is done by creating an 8-bin orientation histogram in each 4 x 4 neighborhood of the 16 x 16 neighborhood around the keypoint



- The obtained 16 histograms are concatenated together to form a single **descriptor** vector of length $4 \times 4 \times 8 = 128$ values
- This descriptor is invariant to scale changes and brightness affine transformations

SIFT: Descriptors Matching

- The final step is to find for each descriptor in one image the corresponding similar descriptor in the other image
- This is done by comparing descriptors across the two images using Euclidean distance
- For each descriptor in one image, the corresponding descriptor in the other image is the one that gives the least Euclidean distance (preferably 0 distance)

Compare the descriptor of this blob with the descriptors of all other blobs in Image 2 and find the one that gives zero (or minimum) distance

From Image 1




Image 2



SIFT: Descriptors Matching

- Finally, we find the average shift between the corresponding blobs across the two images
- This shift is used to adjust how the two images will be stitched

Image 1



Image 2



Stitched Image

